

## ▼ Transfer Learning MNIST

- Train a simple convnet on the MNIST dataset the first 5 digits [0..4].
- Freeze convolutional layers and fine-tune dense layers for the classification of digits [5..9].

### ▼ Import MNIST data and create 2 datasets with one dataset having digits from 0 to 4 and other from 5 to 9

#### ▼ Import the mnist dataset from keras datasets

```
#Importing important modules
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.callbacks import ModelCheckpoint, EarlyStopping
#Installing Tensorboard for Colab
!pip install tensorboardcolab

[?] Requirement already satisfied: tensorboardcolab in /usr/local/lib/python3.6/dist-packages
```

```
from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

#### ▼ Creating two datasets one with digits below 5 and one with 5 and above

```
x_train_lt5 = x_train[y_train < 5]
y_train_lt5 = y_train[y_train < 5]
x_test_lt5 = x_test[y_test < 5]
y_test_lt5 = y_test[y_test < 5]
```

```
set(y_train_lt5)
```

```
[?] {0, 1, 2, 3, 4}
```

```
x_train_gt5 = x_train[y_train >= 5]
y_train_gt5 = y_train[y_train >= 5] - 5 # make classes start at 0 for
x_test_gt5 = x_test[y_test >= 5] # np_utils.to_categorical
y_test_gt5 = y_test[y_test >= 5] - 5
```

```
set(y_train_gt5)
```

```
↳ {0, 1, 2, 3, 4}
```

```
set(y_train)
```

```
↳ {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
```

## ▼ Check

Verify shapes of x\_train, y\_train, x\_test and y\_test for both the datasets with the below given shape

```
x_test.shape
```

```
↳ (10000, 28, 28)
```

```
print(x_train_lt5.shape)
print(y_train_lt5.shape)
print(x_test_lt5.shape)
print(y_test_lt5.shape)
```

```
↳ (30596, 28, 28)
(30596,)
(5139, 28, 28)
(5139,)
```

```
print(x_train_gt5.shape)
print(y_train_gt5.shape)
print(x_test_gt5.shape)
print(y_test_gt5.shape)
```

```
↳ (29404, 28, 28)
(29404,)
(4861, 28, 28)
(4861,)
```

## ▼ Let us take only the dataset (x\_train, y\_train, x\_test, y\_test) for Integers 0 to 4 i.e.

Reshape x\_train and x\_test to a 4 Dimensional array (channel = 1) to pass it into the model

```
x_train_lt5 = x_train_lt5.reshape(x_train_lt5.shape[0], 28, 28, 1)
```

```
x_train_lt5[0].shape
```

```
↳ (28, 28, 1)
```

```
x_test_lt5 = x_test_lt5.reshape(x_test_lt5.shape[0], 28, 28, 1)
```

```
x_test_lt5[0].shape
```

↳ (28, 28, 1)

- ▼ Change into float32 datatype and Normalize x\_train and x\_test by dividing it by

```
x_train_lt5 = x_train_lt5.astype('float32')
x_test_lt5 = x_test_lt5.astype('float32')

#Normalizing the input
x_train_lt5 /= 255.0
x_test_lt5 /= 255.0
```

- ▼ Check

Verify the shapes of the X\_train and X\_test with the shapes given below.

```
print('x_train shape:', x_train_lt5.shape)
print(x_train_lt5.shape[0], 'train samples')
print(x_test_lt5.shape[0], 'test samples')
```

↳ x\_train shape: (30596, 28, 28, 1)  
 30596 train samples  
 5139 test samples

```
print('X_train shape:', x_train_lt5.shape)
print('X_test shape:', x_test_lt5.shape)
```

↳ X\_train shape: (30596, 28, 28, 1)  
 X\_test shape: (5139, 28, 28, 1)

```
batch_size = 128
num_classes = 5
epochs = 5
```

- ▼ Use One-hot encoding to divide y\_train and y\_test into required no of output classes

```
# convert class vectors to binary class matrices
y_train_lt5 = keras.utils.to_categorical(y_train_lt5, num_classes)
y_test_lt5 = keras.utils.to_categorical(y_test_lt5, num_classes)
```

```
set(y_train_lt5)
```

↳ {0, 1, 2, 3, 4}

```
# input image dimensions
img_rows, img_cols = 28, 28
```

```
#Keras expects data to be in the format (N_E,N_H,N_W,N_C)
#N_E = Number of Examples, N_H = height, N_W = Width, N_C = Number of Channels.
x_train_lt5 = x_train_lt5.reshape(x_train_lt5.shape[0], img_rows, img_cols, 1)
x_test_lt5 = x_test_lt5.reshape(x_test_lt5.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)
```

- ▼ Build a sequential model with 2 Convolutional layers with 32 kernels of size (3, size (2,2) followed by a drop out layer to be trained for classification of digits 0

```
#Initialize the model
model = Sequential()

#Add a Convolutional Layer with 32 filters of size 3X3 and activation function as 'ReLU'
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))

#Add a Convolutional Layer with 64 filters of size 3X3 and activation function as 'ReLU'
model.add(Conv2D(64, (3, 3), activation='relu'))

#Add a MaxPooling Layer of size 2X2
model.add(MaxPooling2D(pool_size=(2, 2)))

#Apply Dropout with 0.25 probability
model.add(Dropout(0.25))
```

- ▼ Post that flatten the data and add 2 Dense layers with 128 neurons and neuron 'relu' and 'softmax' respectively. Add dropout layer inbetween if necessary

```
#Flatten the layer
model.add(Flatten())

#Add Fully Connected Layer with 128 units and activation function as 'ReLU'
model.add(Dense(128, activation='relu'))
#Apply Dropout with 0.5 probability
model.add(Dropout(0.5))

#Add Fully Connected Layer with 10 units and activation function as 'softmax'
model.add(Dense(num_classes, activation='softmax'))
```

- ▼ Print the training and test accuracy for 5 epochs

```
from keras.optimizers import Adam
from keras.losses import categorical_crossentropy

#To use adam optimizer for learning weights with learning rate = 0.001
```

```
optimizer = Adam(lr=0.001)
#Set the loss function and optimizer for the model training
model.compile(loss=categorical_crossentropy,
                optimizer=optimizer,
                metrics=['accuracy'])

#Import tensorboardcolab modules for creating a tensorboard call back which will passed in
from tensorboardcolab import TensorBoardColab, TensorBoardColabCallback

#Tensorboard callback is going to be added to model.fit function to draw graphs of loss va
tbc = TensorBoardColab()

↳ Wait for 8 seconds...
TensorBoard link:
https://c5246f93.ngrok.io

#Adding Early stopping callback to the fit function is going to stop the training,
#if the val_loss is not going to change even '0.001' for more than 10 continous epochs

early_stopping = EarlyStopping(monitor='val_loss', min_delta=0.001, patience=10)

#Adding Model Checkpoint callback to the fit function is going to save the weights whenever
#Hence saving the best weights occurred during training

model_checkpoint = ModelCheckpoint('mnist_cnn_checkpoint_{epoch:02d}_loss{val_loss:.4f}.h
                                    monitor='val_loss',
                                    verbose=1,
                                    save_best_only=True,
                                    save_weights_only=True,
                                    mode='auto',
                                    period=1)

#Training on the dataset and adding the all the callbacks to the fit function.
#Once the training starts, results start appearing on Tensorboard after 1 epoch
model.fit(x_train_lt5, y_train_lt5,
            batch_size=batch_size,
            epochs=epochs,
            verbose=1,
            validation_data=(x_test_lt5, y_test_lt5),
            callbacks=[TensorBoardColabCallback(tbc),early_stopping,model_checkpoint])
```

↳

```
Train on 30596 samples, validate on 5139 samples
Epoch 1/5
30596/30596 [=====] - 5s 153us/step - loss: 0.1168 - acc: 0.

Epoch 00001: val_loss improved from inf to 0.01987, saving model to mnist_cnn_checkpo
Epoch 2/5
30596/30596 [=====] - 4s 142us/step - loss: 0.0318 - acc: 0.

Epoch 00002: val_loss improved from 0.01987 to 0.01163, saving model to mnist_cnn_che
Epoch 3/5
30596/30596 [=====] - 4s 144us/step - loss: 0.0220 - acc: 0.

Epoch 00003: val_loss improved from 0.01163 to 0.01009, saving model to mnist_cnn_che
Epoch 4/5
30596/30596 [=====] - 4s 140us/step - loss: 0.0161 - acc: 0.

Epoch 00004: val_loss improved from 0.01009 to 0.00533, saving model to mnist_cnn_che
Epoch 5/5
30596/30596 [=====] - 4s 143us/step - loss: 0.0130 - acc: 0.

Epoch 00005: val_loss improved from 0.00533 to 0.00438, saving model to mnist_cnn_che
<keras.callbacks.History at 0x7f7030112828>
```

```
#Testing the model on test set
score = model.evaluate(x_test_lt5, y_test_lt5)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

→ 5139/5139 [=====] - 1s 109us/step  
 Test loss: 0.004378508583054539  
 Test accuracy: 0.9986378672893559

- ▼ Use the model trained on 0 to 4 digit classification and train it on the dataset w
- ▼ Transfer learning keeping only the dense layers to be trainable)
- Make only the dense layers to be trainable and convolutional layers to be non-t
- ▼ Check model summary to see model layer names

```
for layers in model.layers:
    print(layers.name)
    if('dense' not in layers.name):
        layers.trainable = False
        print(layers.name + 'is not trainable\n')
    if('dense' in layers.name):
        print(layers.name + ' is trainable\n')
```

→

```
conv2d_3
conv2d_3is not trainable

conv2d_4
conv2d_4is not trainable

max_pooling2d_2
max_pooling2d_2is not trainable

dropout_3
dropout_3is not trainable

flatten_2
flatten_2is not trainable

dense_3
dense_3 is trainable

dropout_4
dropout_4is not trainable

dense_4
dense_4 is trainable
```

Do the required preprocessing for `x_train_gt5` also same as `x_train_lt5` and `y_train_lt5`

1. Reshape
2. Change to float32 datatype
3. Normalize (dividing with 255)
4. `y_train` and `y_test` Convert into one-hot vectors

### Reshape

```
x_train_gt5 = x_train_gt5.reshape(x_train_gt5.shape[0],28,28,1)
```

```
x_train_gt5.shape
```

```
↳ (29404, 28, 28, 1)
```

```
set(y_train_gt5)
```

```
↳ {0, 1, 2, 3, 4}
```

```
x_test_gt5 = x_test_gt5.reshape(x_test_gt5.shape[0],28,28,1)
```

Change to Float and Normalize

```
x_train_gt5 = x_train_gt5.astype('float32')
x_test_gt5 = x_test_gt5.astype('float32')
```

```
#Normalizing the input
x_train_gt5 /= 255.0
x_test_gt5 /= 255.0
```

## OneHot

```
# convert class vectors to binary class matrices
y_train_gt5 = keras.utils.to_categorical(y_train_gt5, num_classes)
y_test_gt5 = keras.utils.to_categorical(y_test_gt5, num_classes)
```

### ▼ Check

Verify the shapes with the given below.

```
print(x_train_gt5.shape)
print(y_train_gt5.shape)
print(x_test_gt5.shape)
print(y_test_gt5.shape)
```

```
⇒ (29404, 28, 28, 1)
(29404, 5)
(4861, 28, 28, 1)
(4861, 5)
```

### ▼ Print the accuracy for classification of digits 5 to 9

```
#Training on the dataset and adding the all the callbacks to the fit function.
#Once the training starts, results start appearing on Tensorboard after 1 epoch
model.fit(x_train_gt5, y_train_gt5,
           batch_size=batch_size,
           epochs=epochs,
           verbose=1,
           validation_data=(x_test_gt5, y_test_gt5),
           callbacks=[TensorBoardColabCallback(tbc), early_stopping, model_checkpoint])
```

```
⇒
```

```

Train on 29404 samples, validate on 4861 samples
Epoch 1/5
 128/29404 [...........................] - ETA: 5s - loss: 5.6763 - acc: 0.2812/u
'Discrepancy between trainable weights and collected trainable'
29404/29404 [=====] - 4s 147us/step - loss: 0.1961 - acc: 0.

Epoch 00001: val_loss did not improve from 0.00438
Epoch 2/5
29404/29404 [=====] - 4s 141us/step - loss: 0.0550 - acc: 0.

Epoch 00002: val_loss did not improve from 0.00438
Epoch 3/5
29404/29404 [=====] - 4s 142us/step - loss: 0.0395 - acc: 0.

Epoch 00003: val_loss did not improve from 0.00438
Epoch 4/5
29404/29404 [=====] - 4s 143us/step - loss: 0.0296 - acc: 0.

Epoch 00004: val_loss did not improve from 0.00438
Epoch 5/5
29404/29404 [=====] - 4s 144us/step - loss: 0.0245 - acc: 0.

Epoch 00005: val_loss did not improve from 0.00438
<keras.callbacks.History at 0x7f703d121f60>

```

```

#Testing the model on test set
score = model.evaluate(x_test_gt5, y_test_gt5)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

⇒ 4861/4861 [=====] - 1s 109us/step
Test loss: 0.020349998070301942
Test accuracy: 0.9942398683398478

```

## ▼ Text classification using TF-IDF

### ▼ Load the dataset from sklearn.datasets

```

from sklearn.datasets import fetch_20newsgroups

categories = ['alt.atheism', 'soc.religion.christian', 'comp.graphics', 'sci.med']

```

### ▼ Training data

```

twenty_train = fetch_20newsgroups(subset='train', categories=categories, shuffle=True, ran

```

```

⇒

```

Downloading 20news dataset. This may take a few minutes.

## ▼ Test data

```
twenty_test = fetch_20newsgroups(subset='test', categories=categories, shuffle=True, rando
```

- ▼ a. You can access the values for the target variable using .target attribute
- b. You can access the name of the class in the target variable with .target\_nam

```
twenty_train.target
```

```
↳ array([1, 1, 3, ..., 2, 2, 2])
```

```
twenty_train
```

```
↳
```

```
'From: dfuller@portal.hq.videocart.com (Dave Fuller)\nSubject: Re: thoughts on chri
'From: halsall@murray.fordham.edu (Paul Halsall)\nSubject: Bible Unsuitable for New
'From: twain@carson.u.washington.edu (Barbara Hlavin)\nSubject: Re: Is MSG sensitiv
'From: Desiree\_Bradley@mindlink.bc.ca (Desiree Bradley)\nSubject: Doing the work of
'From: db7n+@andrew.cmu.edu (D. Andrew Byler)\nSubject: Re: What WAS the immaculate
'From: Donald Mackie <Donald\_Mackie@med.umich.edu>\nSubject: Re: Seeking advice/exp
"From: weilej@cary115.its.rpi.edu (Jason Lee Weiler)\nSubject: Re: need a viewer fo
"From: ednobles@sacam.OREN.ORTN.EDU (Edward d Nobles)\nSubject: POV .TGA's and Spee
"Subject: .GL and .FLI specs\nFrom: arthur@qedbbs.com (Arthur Choung)\nOrganization
'From: jbrown@batman.bmd.trw.com\nSubject: Re: Death Penalty / Gulf War (long)\nLin
"From: bobbe@vice.ICO.TEK.COM (Robert Beauchaine)\nSubject: Re: thoughts on christi
"From: rcomg@melomys.co.rmit.oz.AU (Mark Gregory)\nSubject: AVI file format?\nSumma
'From: jimh@carson.u.washington.edu (James Hogan)\nSubject: Re: Yet more Rushdie [R
'From: alastair@farli.otago.ac.nz (Alastair Thomson)\nSubject: Does \'Just/justifia
"From: will@futon.webo.dg.com (Will Taber)\nSubject: [soc.religion.christian] Re: T
'From: mathew <mathew@mantis.co.uk>\nSubject: Alt.Atheism FAQ: Introduction to Athe
'From: timmbake@mcl.ucsb.edu (Bake Timmons)\nSubject: Re: Amusing atheists and agno
'From: kmr4@po.CWRU.edu (Keith M. Ryan)\nSubject: Re: The Inimitable Rushdie\nOrgan
"From: tonyo@pendragon.CNA.TEK.COM (Tony Ozrelic)\nSubject: Need info on cc:Mail fi
"From: scornd7@technet.sg (Tang Chang Thai)\nSubject: Re: InterViews graphics packa
'From: mmatusev@radford.vak12ed.edu (Melissa N. Matusevich)\nSubject: Re: Emphysema
'From: livesey@solntze.wpd.sgi.com (Jon Livesey)\nSubject: Re: <Political Atheists?
'From: rjs2@po.cwru.edu (Richard J. Szanto)\nSubject: Re: When are two people marri
"From: rws2v@uvacs.cs.Virginia.EDU (Richard Stoakley)\nSubject: Need a good concave
'From: Geoffrey Hansen@mindlink.bc.ca (Geoffrey Hansen)\nSubject: Re: VESA on the S
'From: erh0362@tesla.njit.edu\nSubject: Mormon beliefs about bastards\nOrganization
'From: bil@okcforum.osrhe.edu (Bill Conner)\nSubject: Re: some thoughts.\nNntp-Post
"From: mary@uicsl.cs1.uiuc.edu (Mary E. Allison)\nSubject: Re: Is MSG sensitivity s
'From: sieferme@stein.u.washington.edu (Eric Sieberman)\nSubject: Re: some thoughts
'From: Lars.Jorgensen@p7.syntax.bbs.bad.se (Lars Jorgensen)\nSubject: Externel proc
'From: maridai@comm.mot.com (Marida Ignacio)\nSubject: Re: "Accepting Jesus in your
'From: h8902939@hkuxa.hku.hk (Abel)\nSubject: Developable Surface\nNntp-Posting-Hos
'From: wsun@jeeves.ucsd.edu (Fiberman)\nSubject: Re: Is MSG sensitivity superstitio
'From: mcovingt@aisun3.ai.uga.edu (Michael Covington)\nSubject: Re: When are two pe
'From: Patrick C Leger <p11u+@andrew.cmu.edu>\nSubject: Re: thoughts on christians\
"From: diablo.UUCP!cboesel (Charles Boesel)\nSubject: Re: Postscript drawing prog\n
"From: I3150101@dbstu1.rz.tu-bs.de (Benedikt Rosenau)\nSubject: Re: The Inimitable
"From: nichael@bbn.com (Nichael Cramer)\nSubject: Re: Dead Sea Scrolls\nReply-To: nichael@bbn.com
'From: atterlep@vela.acs.oakland.edu (Cardinal Ximenez)\nSubject: Re: A question th
'From: cobb@alexia.lis.uiuc.edu (Mike Cobb)\nSubject: Re: After 2000 years, can we
'From: parkin@Eng.Sun.COM (Michael Parkin)\nSubject: Re: Being right about messiahs
'From: kxgst1+pitt.edu (Kenneth Gilbert)\nSubject: Re: Any info. on Vasomotor Rhin
"From: labson@borneo.corp.sgi.com (Joel Labson)\nSubject: Maybe?????\nOrganization:
'From: dyer@spdcc.com (Steve Dyer)\nSubject: Re: Thrush ((was: Good Grief! (was Re:
...],  

'filenames': array(['/root/scikit_learn_data/20news_home/20news-bydate-train/comp.gr
        '/root/scikit_learn_data/20news_home/20news-bydate-train/comp.graphics/38479'
        '/root/scikit_learn_data/20news_home/20news-bydate-train/soc.religion.christi
        ...,
        '/root/scikit_learn_data/20news_home/20news-bydate-train/sci.med/58112',
        '/root/scikit_learn_data/20news_home/20news-bydate-train/sci.med/58578',
        '/root/scikit_learn_data/20news_home/20news-bydate-train/sci.med/58895'],
        dtype='|<U86'),  

'target': array([1, 1, 3, ..., 2, 2, 2]),  

'target_names': ['alt.atheism',
 'comp.graphics',
 'sci.med',
 'soc.religion.christian']}
```







